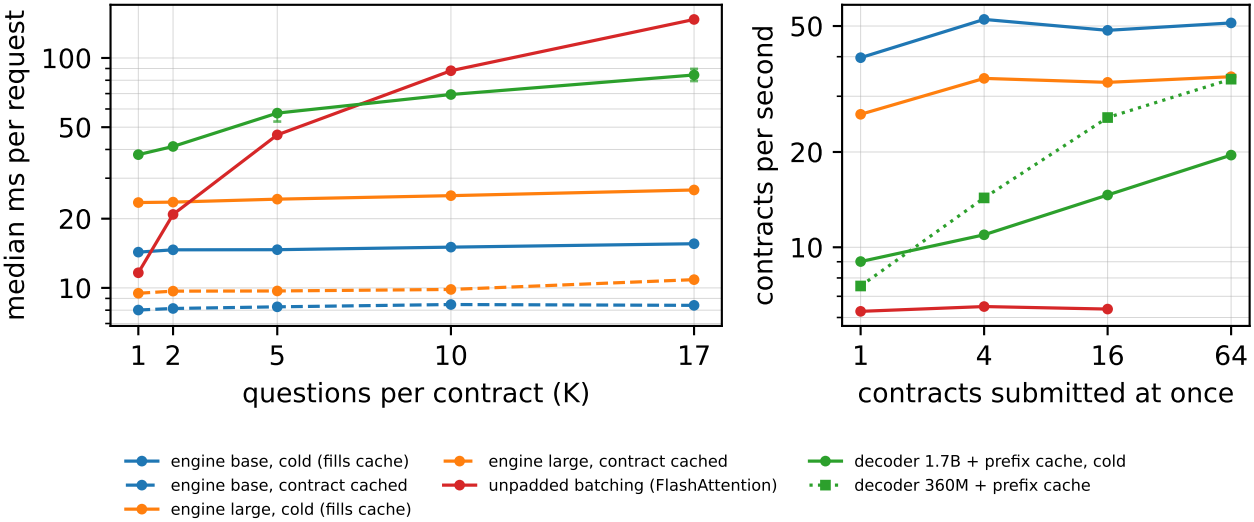
median ms per request
100
50
20
10
1
2
5
10
17
questions per contract (K)
contracts per second
50
20
10
1
4
16
64
contracts submitted at once
engine base, cold (fills cache)
engine base, contract cached
engine large, cold (fills cache)
engine large, contract cached
unpadded batching (FlashAttention)
decoder 1.7B + prefix cache, cold
decoder 360M + prefix cache