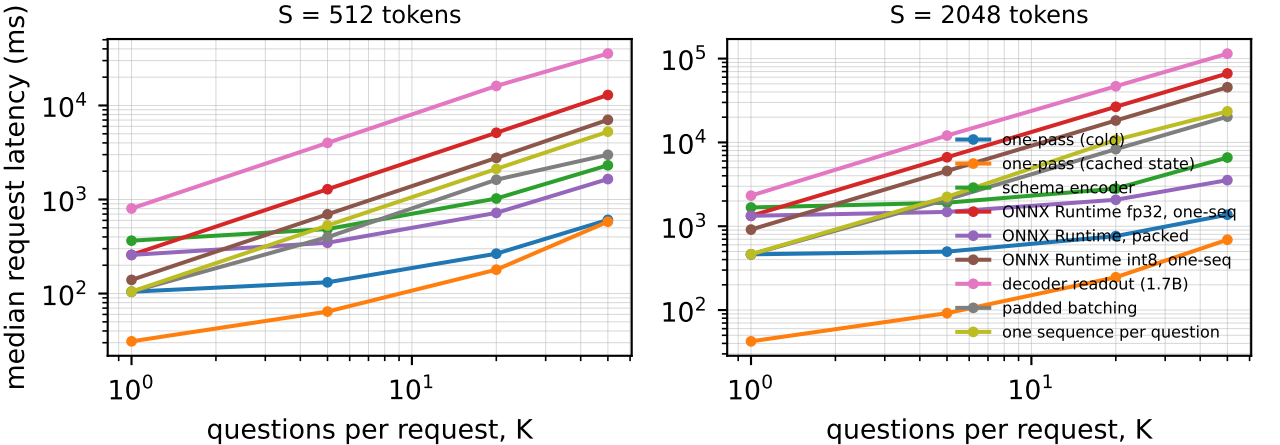
S = 512 tokens
S = 2048 tokens
median request latency (ms)
questions per request, K
questions per request, K
one-pass (cold)
one-pass (cached state)
schema encoder
ONNX Runtime fp32, one-seq
ONNX Runtime, packed
ONNX Runtime int8, one-seq
decoder readout (1.7B)
padded batching
one sequence per question