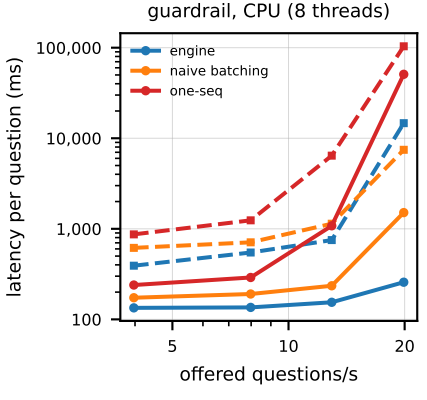
guardrail, CPU (8 threads)
latency per question (ms)
100,000
10,000
1,000
100
5
10
20
offered questions/s
engine
naive batching
one-seq

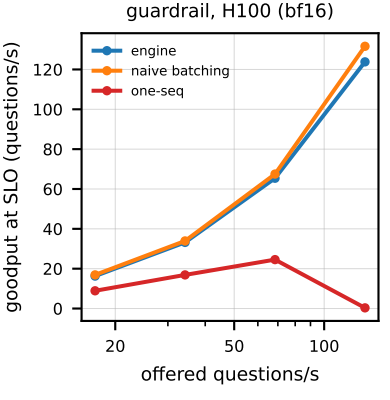
guardrail, H100 (bf16)
goodput at SLO (questions/s)
120
100
80
60
40
20
0
20
50
100
offered questions/s
engine
naive batching
one-seq

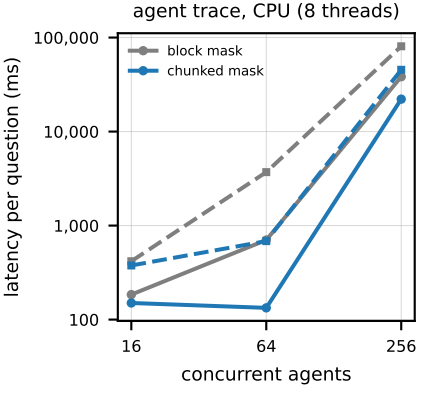
agent trace, CPU (8 threads)
latency per question (ms)
100,000
10,000
1,000
100
16
64
256
concurrent agents
block mask
chunked mask